

# OLA - Ensemble Learning

## Problem Statement

Recruiting and retaining drivers is seen by industry watchers as a tough battle for Ola. Churn among drivers is high and it's very easy for drivers to stop working for the service on the fly or jump to Uber depending on the rates.

As the companies get bigger, the high churn could become a bigger problem. To find new drivers, Ola is casting a wide net, including people who don't have cars for jobs. But this acquisition is really costly. Losing drivers frequently impacts the morale of the organization and acquiring new drivers is more expensive than retaining existing ones.

You are working as a data scientist with the Analytics Department of Ola, focused on driver team attrition. You are provided with the monthly information for a segment of drivers for 2019 and 2020 and tasked to predict whether a driver will be leaving the company or not based on their attributes like

Demographics (city, age, gender etc.)

- Tenure information (joining date, Last Date)
- Historical data regarding the performance of the driver (Quarterly rating, Monthly business acquired, grade, Income)

## Column Profiling:

- MMMM-YY : Reporting Date (Monthly)
- Driver\_ID : Unique id for drivers
- Age : Age of the driver
- Gender : Gender of the driver – Male : 0, Female: 1
- City : City Code of the driver
- Education\_Level : Education level – 0 for 10+ ,1 for 12+ ,2 for graduate
- Income : Monthly average Income of the driver
- Date Of Joining : Joining date for the driver
- LastWorkingDate : Last date of working for the driver
- Joining Designation : Designation of the driver at the time of joining
- Grade : Grade of the driver at the time of reporting
- Total Business Value : The total business value acquired by the driver in a month (negative business indicates cancellation/refund or car EMI adjustments)
- Quarterly Rating : Quarterly rating of the driver: 1,2,3,4,5 (higher is better)

```
In [1]: # Analysis
        #Importing Required Libraries

import numpy as np
```

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import statsmodels.api as sm
from datetime import datetime
```

```
In [2]: df=pd.read_csv('ola_driver_scaler.csv')
df.head()
```

```
Out[2]:
```

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	LastWorkingDate	Joining Designation	Grade	Total Business Value	Quarterly Rating
0	0	01/01/19	1	28.0	0.0	C23	2	57387	24/12/18					
1	1	02/01/19	1	28.0	0.0	C23	2	57387	24/12/18					
2	2	03/01/19	1	28.0	0.0	C23	2	57387	24/12/18					
3	3	11/01/20	2	31.0	0.0	C7	2	67016	11/06/20					
4	4	12/01/20	2	31.0	0.0	C7	2	67016	11/06/20					

```
In [3]: df.shape
```

```
Out[3]: (19104, 14)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19104 entries, 0 to 19103
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Unnamed: 0                            19104 non-null  int64
 1   MMM-YY                                19104 non-null  object
 2   Driver_ID                             19104 non-null  int64
 3   Age                                    19043 non-null  float64
 4   Gender                                19052 non-null  float64
 5   City                                  19104 non-null  object
 6   Education_Level                       19104 non-null  int64
 7   Income                                19104 non-null  int64
 8   Dateofjoining                         19104 non-null  object
 9   LastWorkingDate                       1616 non-null   object
10   Joining Designation                   19104 non-null  int64
11   Grade                                 19104 non-null  int64
12   Total Business Value                  19104 non-null  int64
13   Quarterly Rating                      19104 non-null  int64
dtypes: float64(2), int64(8), object(4)
memory usage: 2.0+ MB
```

```
In [5]: #calculating Total missing values
df.isnull().sum()
```

```
Out[5]: Unnamed: 0          0
      MMM-YY          0
      Driver_ID        0
      Age             61
      Gender          52
      City            0
      Education_Level  0
      Income           0
      Dateofjoining    0
      LastWorkingDate  17488
      Joining Designation 0
      Grade            0
      Total Business Value 0
      Quarterly Rating  0
      dtype: int64
```

## Observations:

- There are some Missing Values
- Lets impute them with KNN Imputation

```
In [7]: #Dropping Unnamed column which is not Required
df.drop(['Unnamed: 0'], axis=1, inplace=True)
```

```
In [8]: #Converting series to numpy array to impute KNN Imputation
age=df['Age']
age=age.to_numpy(dtype='float32').reshape(-1,1)
gender=df['Gender']
gender=gender.to_numpy(dtype='float32').reshape(-1,1)
```

```
In [9]: #Imputing Missing values with KNN Imoutation
from sklearn.impute import KNNImputer
imputer = KNNImputer(n_neighbors=5)

df['Age']=imputer.fit_transform(age)
df['Gender']=imputer.fit_transform(gender)
```

```
In [10]: #After applying KNN imputaion it converted to numpy array so converted to Series
df['Age'] = pd.Series(df['Age'])
df['Gender'] = pd.Series(df['Gender'])
```

```
In [11]: #converting float to INT type
df['Age'] = df['Age'].astype('int')
df['Gender'] = df['Gender'].astype('int')
```

```
In [12]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19104 entries, 0 to 19103
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   MMM-YY                                19104 non-null  object
1   Driver_ID                            19104 non-null  int64
2   Age                                  19104 non-null  int32
3   Gender                              19104 non-null  int32
4   City                                 19104 non-null  object
5   Education_Level                      19104 non-null  int64
6   Income                              19104 non-null  int64
7   Dateofjoining                       19104 non-null  object
8   LastWorkingDate                     1616 non-null   object
9   Joining Designation                 19104 non-null  int64
10  Grade                               19104 non-null  int64
11  Total Business Value                19104 non-null  int64
12  Quarterly Rating                    19104 non-null  int64
dtypes: int32(2), int64(7), object(4)
memory usage: 1.7+ MB

```

```

In [13]: #converting columes to datetime format to subtract them
df['MMM-YY'] = pd.to_datetime(df['MMM-YY'])
df['Dateofjoining'] = pd.to_datetime(df['Dateofjoining'])
df['LastWorkingDate'] = pd.to_datetime(df['LastWorkingDate'])

```

```

In [14]: df_new=df.groupby(['Driver_ID','Quarterly Rating']).agg({'MMM-YY':'last','Age':'last',
                                                                    'Education_Level':'last',
                                                                    'LastWorkingDate':'last',
                                                                    'Grade':'first',
                                                                    'Total Business Value':'last'})

df_new

```

Out[14]:

		MMM-YY	Age	Gender	City	Education_Level	Dateofjoining	LastWorkingDate
Driver_ID	Quarterly Rating							
1	2	2019-03-01	28	0	C23	2	2018-12-24	2019-03-11
2	1	2020-12-01	31	0	C7	2	2020-11-06	Na1
4	1	2020-04-01	43	0	C13	2	2019-12-07	2020-04-27
5	1	2019-03-01	29	0	C9	0	2019-01-09	2019-03-07
6	1	2020-09-01	31	1	C11	1	2020-07-31	Na1
...	...	...	...	...	...	...	...	...
2787	1	2019-06-01	28	1	C20	2	2018-07-21	2019-06-20
	2	2019-03-01	28	1	C20	2	2018-07-21	Na1
2788	1	2020-06-01	29	0	C27	2	2020-06-08	Na1
	2	2020-12-01	30	0	C27	2	2020-06-08	Na1
	3	2020-09-01	30	0	C27	2	2020-06-08	Na1

4023 rows × 11 columns



```
In [15]: #Not all columns are present so applying reset index
df1=df_new.reset_index()
```

```
In [16]: # sorted values to find diff in driver rating
df1.sort_values(by=['Driver_ID', 'MMM-YY'], inplace=True)
```

```
In [17]: df1['Rating_Increased']=df1.groupby(['Driver_ID'])['Quarterly Rating'].diff()
df1['Rating_Decreased']=df1.groupby(['Driver_ID'])['Quarterly Rating'].diff()
```

```
In [18]: # Created new Columns based on driver rating
df1['Rating_Increased']=df1['Rating_Increased'].apply(lambda x: 1 if x>0 else 0 )
df1['Rating_Decreased']=df1['Rating_Decreased'].apply(lambda x: 1 if x<0 else 0 )
```

```
In [19]: df_new=df1.groupby(['Driver_ID']).agg({'MMM-YY':'last', 'Age':'last', 'Gender':'last',
'Dateofjoining':'first', 'LastWorkingDate':'last', 'Education_Level':'last',
'Grade':'last', 'Income':'last', 'Total Business':'last',
'Rating_Decreased':'last', 'Quarterly Rating':'last'})
df_new
```

Out[19]:

	MMM- YY	Age	Gender	City	Education_Level	Dateofjoining	LastWorkingDate	Joini Designati
Driver_ID								
1	2019-03-01	28	0	C23	2	2018-12-24	2019-03-11	
2	2020-12-01	31	0	C7	2	2020-11-06	NaT	
4	2020-04-01	43	0	C13	2	2019-12-07	2020-04-27	
5	2019-03-01	29	0	C9	0	2019-01-09	2019-03-07	
6	2020-12-01	31	1	C11	1	2020-07-31	NaT	
...	...	...	...	...	...	...	...	...
2784	2020-12-01	34	0	C24	0	2015-10-15	NaT	
2785	2020-10-01	34	1	C9	0	2020-08-28	2020-10-28	
2786	2019-09-01	45	0	C19	0	2018-07-31	2019-09-22	
2787	2019-06-01	28	1	C20	2	2018-07-21	2019-06-20	
2788	2020-12-01	30	0	C27	2	2020-06-08	NaT	

2381 rows × 14 columns



```
In [20]: #Not all columns are present so applying reset index
df=df_new.reset_index()
df
```

Out[20]:

	Driver_ID	MMM-YY	Age	Gender	City	Education_Level	Dateofjoining	LastWorkingDate	Dt
0	1	2019-03-01	28	0	C23	2	2018-12-24	2019-03-11	
1	2	2020-12-01	31	0	C7	2	2020-11-06	NaT	
2	4	2020-04-01	43	0	C13	2	2019-12-07	2020-04-27	
3	5	2019-03-01	29	0	C9	0	2019-01-09	2019-03-07	
4	6	2020-12-01	31	1	C11	1	2020-07-31	NaT	
...	...	...	...	...	...	...	...	...	...
2376	2784	2020-12-01	34	0	C24	0	2015-10-15	NaT	
2377	2785	2020-10-01	34	1	C9	0	2020-08-28	2020-10-28	
2378	2786	2019-09-01	45	0	C19	0	2018-07-31	2019-09-22	
2379	2787	2019-06-01	28	1	C20	2	2018-07-21	2019-06-20	
2380	2788	2020-12-01	30	0	C27	2	2020-06-08	NaT	

2381 rows × 15 columns

```
In [38]: #creation of Target Variable
df['churn']=df['LastWorkingDate']
df['churn'].fillna(0, inplace = True)
df['churn']=df['churn'].apply(lambda x: 0 if x==0 else 1)
```

## deriving basic insights from derived dataset

```
In [22]: df['Rating_Increased'].value_counts()
```

```
Out[22]: 0    2013
         1     368
         Name: Rating_Increased, dtype: int64
```

```
In [23]: df['Rating_Decreased'].value_counts()
```

```
Out[23]: 0    1645
         1     736
         Name: Rating_Decreased, dtype: int64
```

```
In [24]: df['Gender'].value_counts()
```

```
Out[24]: 0    1408
         1     973
         Name: Gender, dtype: int64
```

```
In [25]: df['Education_Level'].value_counts()
```

```
Out[25]: 2     802
         1     795
         0     784
         Name: Education_Level, dtype: int64
```

```
In [39]: df['churn'].value_counts()
```

```
Out[39]: 1     1616
         0      765
         Name: churn, dtype: int64
```

```
In [27]: # Dupluicate value check
if df.shape[0] == df.drop_duplicates().shape[0] :
    print('No duplicates Found')
else:
    duplicates = df.shape[0] - df.drop_duplicates().shape[0]
    print('{} duplicates found'.format(duplicates))
```

No duplicates Found

```
In [28]: df.describe().T
```

```
Out[28]:
```

	count	mean	std	min	25%	50%	75%
<b>Driver_ID</b>	2381.0	1.397559e+03	8.061616e+02	1.0	695.0	1400.0	2100.0
<b>Age</b>	2381.0	3.368459e+01	5.968409e+00	21.0	29.0	33.0	37.0
<b>Gender</b>	2381.0	4.086518e-01	4.916880e-01	0.0	0.0	0.0	1.0
<b>Education_Level</b>	2381.0	1.007560e+00	8.162900e-01	0.0	0.0	1.0	2.0
<b>Joining Designation</b>	2381.0	1.820244e+00	8.414334e-01	1.0	1.0	2.0	2.0
<b>Grade</b>	2381.0	2.082738e+00	9.359084e-01	1.0	1.0	2.0	3.0
<b>Income</b>	2381.0	5.933416e+04	2.838367e+04	10747.0	39104.0	55315.0	75986.0
<b>Total Business Value</b>	2381.0	4.586742e+06	9.127115e+06	-1385530.0	0.0	817680.0	4173650.0
<b>Rating_Increased</b>	2381.0	1.545569e-01	3.615577e-01	0.0	0.0	0.0	0.0
<b>Rating_Decreased</b>	2381.0	3.091138e-01	4.622253e-01	0.0	0.0	0.0	1.0
<b>Quarterly Rating</b>	2381.0	1.427971e+00	8.098389e-01	1.0	1.0	1.0	2.0
<b>churn</b>	2381.0	1.000000e+00	0.000000e+00	1.0	1.0	1.0	1.0

## Observations:

- Some Missing Values are present in Given Dataset, Imputed with KNN Imputation.
- No Duplicates are found in Dataset after performing Feature engineering.
- No change in monthly income so not including that column for further analysis.
- Major Drivers are male according to Dataset.



- Total of 368 drivers rating has been increased.
- Total of 973 drivers rating has been Decreased.
- Major of Drivers are Graduated according to Dataset.
- There is a imbalance between churn and non churn data we need to use SMOTE techniques later to create Balance.

```
In [29]: df_numerical=df.select_dtypes(exclude='object')
df_categorical=df.select_dtypes(include='object')
```

```
In [30]: df_numerical.columns
```

```
Out[30]: Index(['Driver_ID', 'MMM-YY', 'Age', 'Gender', 'Education_Level',
              'Dateofjoining', 'Joining Designation', 'Grade', 'Income',
              'Total Business Value', 'Rating_Increased', 'Rating_Decreased',
              'Quarterly Rating', 'churn'],
              dtype='object')
```

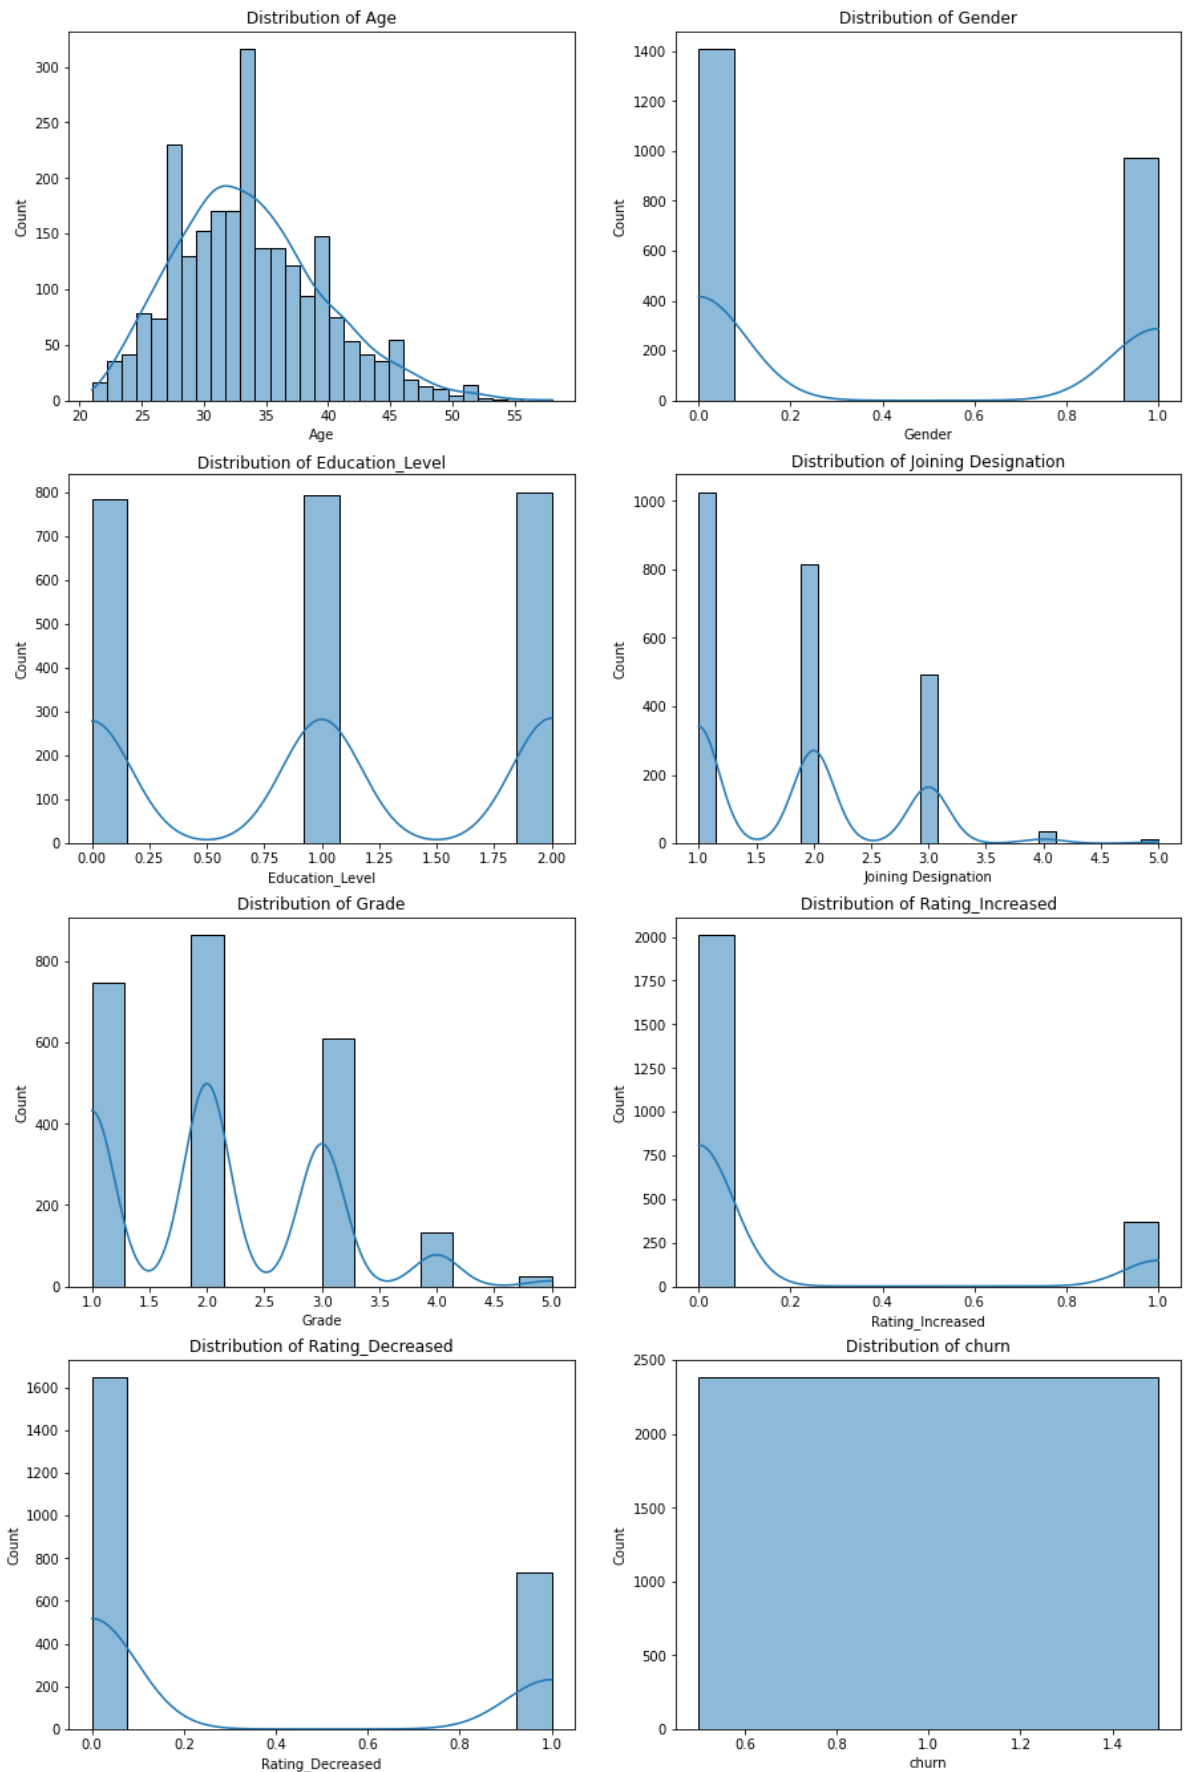
```
In [31]: df_categorical.columns
```

```
Out[31]: Index(['City', 'LastWorkingDate'], dtype='object')
```

## Univariate Analysis

```
In [32]: fig, axis = plt.subplots(nrows=4, ncols=2, figsize=(15, 10))
fig.subplots_adjust(top=1.9)







sns.histplot(data=df, x="Age", kde=True, ax=axis[0,0])
axis[0,0].set_title('Distribution of Age')
sns.histplot(data=df, x="Gender", kde=True, ax=axis[0,1])
axis[0,1].set_title('Distribution of Gender')
sns.histplot(data=df, x="Education_Level", kde=True, ax=axis[1,0])
axis[1,0].set_title('Distribution of Education_Level')
sns.histplot(data=df, x="Joining Designation", kde=True, ax=axis[1,1])
axis[1,1].set_title('Distribution of Joining Designation')
sns.histplot(data=df, x="Grade", kde=True, ax=axis[2,0])
axis[2,0].set_title('Distribution of Grade')
sns.histplot(data=df, x="Rating_Increased", kde=True, ax=axis[2,1])
axis[2,1].set_title('Distribution of Rating_Increased')
sns.histplot(data=df, x="Rating_Decreased", kde=True, ax=axis[3,0])
axis[3,0].set_title('Distribution of Rating_Decreased')
sns.histplot(data=df, x="churn", kde=True, ax=axis[3,1])
axis[3,1].set_title('Distribution of churn')
plt.show()
```



```
In [33]: fig, axis = plt.subplots(nrows=3, ncols=2, figsize=(12, 10))
fig.subplots_adjust(top=1.9)

sns.boxplot(data=df, x="Age", ax=axis[0,0])
sns.boxplot(data=df, x="Total Business Value", ax=axis[0,1])
sns.boxplot(data=df, x="Education_Level", ax=axis[1,0])
sns.boxplot(data=df, x="Joining Designation", ax=axis[1,1])
sns.boxplot(data=df, x="Grade", ax=axis[2,0])
```


```
sns.boxplot(data=df, x="Income", ax=axis[2,1])  
plt.show()
```



```
In [34]: fig, axis = plt.subplots(nrows=1, figsize=(16, 6))
fig.subplots_adjust(top=1.9)

sns.countplot(data=df, x="City")
```

```
Out[34]: <AxesSubplot:xlabel='City', ylabel='count'>
```



## Observations:

- More drivers are from city c20.
- The average age for Drivers are 30-35.
- Total of 368 drivers rating has been increased.
- Total of 973 drivers rating has been Decreased.
- More Male Drivers are present.

## Bivariate and Multivariate Analysis

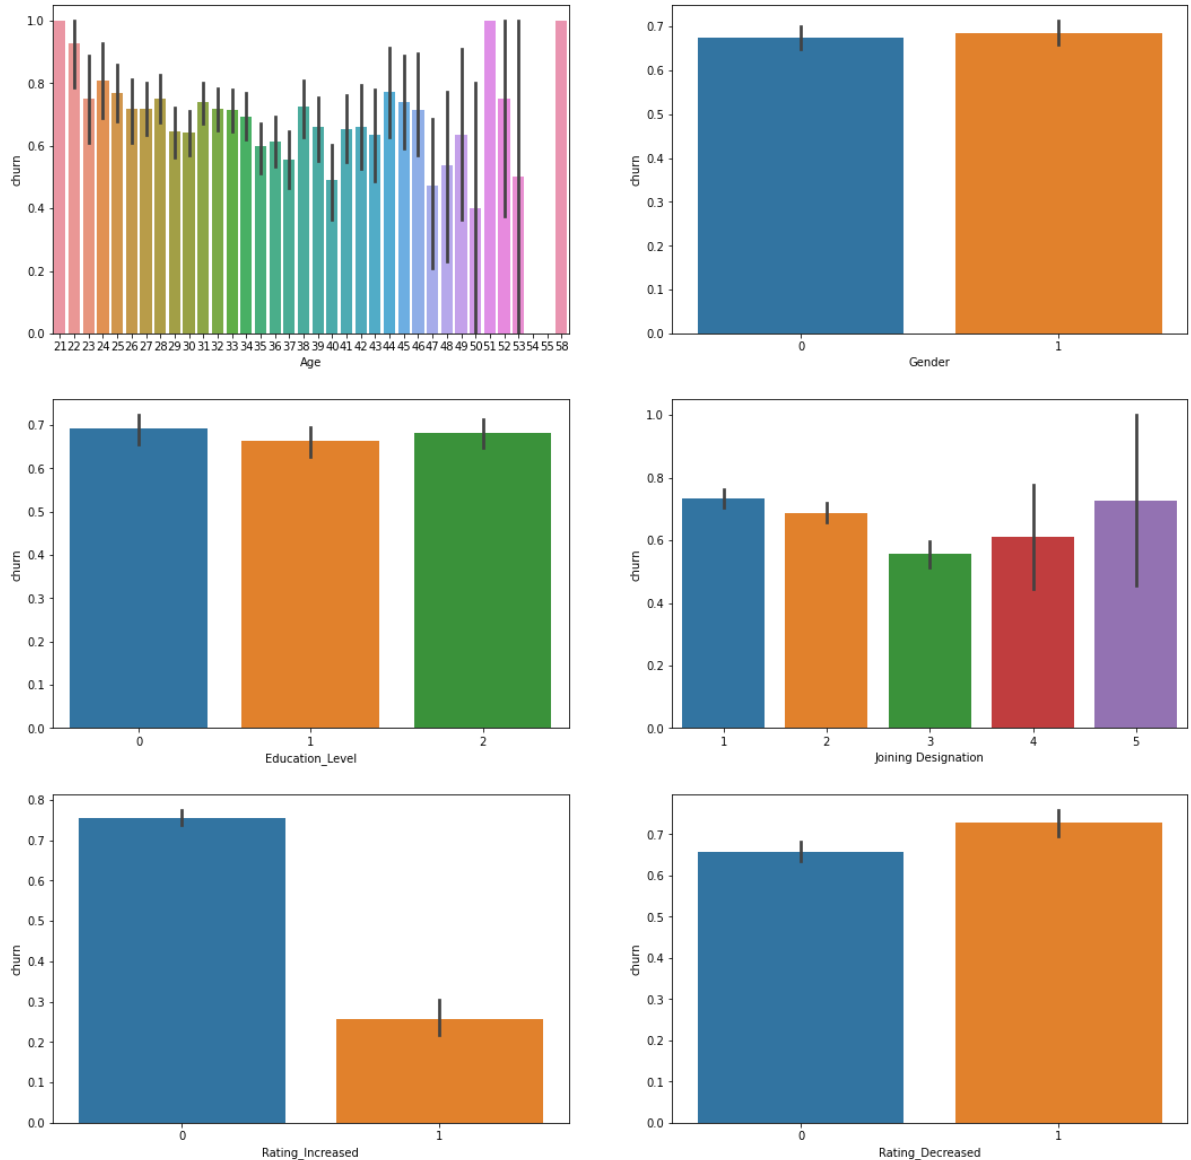
```
In [40]: fig, axis = plt.subplots(nrows=3, ncols=2, figsize=(18, 10))
fig.subplots_adjust(top=1.5)
```

```

sns.barplot(x='Age', y='churn', data=df, ax=axis[0,0])
sns.barplot(x='Gender', y='churn', data=df, ax=axis[0,1])
sns.barplot(x='Education_Level', y='churn', data=df, ax=axis[1,0])
sns.barplot(x='Joining Designation', y='churn', data=df, ax=axis[1,1])
sns.barplot(x='Rating_Increased', y='churn', data=df, ax=axis[2,0])
sns.barplot(x='Rating_Decreased', y='churn', data=df, ax=axis[2,1])

```

Out[40]: <AxesSubplot:xlabel='Rating\_Decreased', ylabel='churn'>



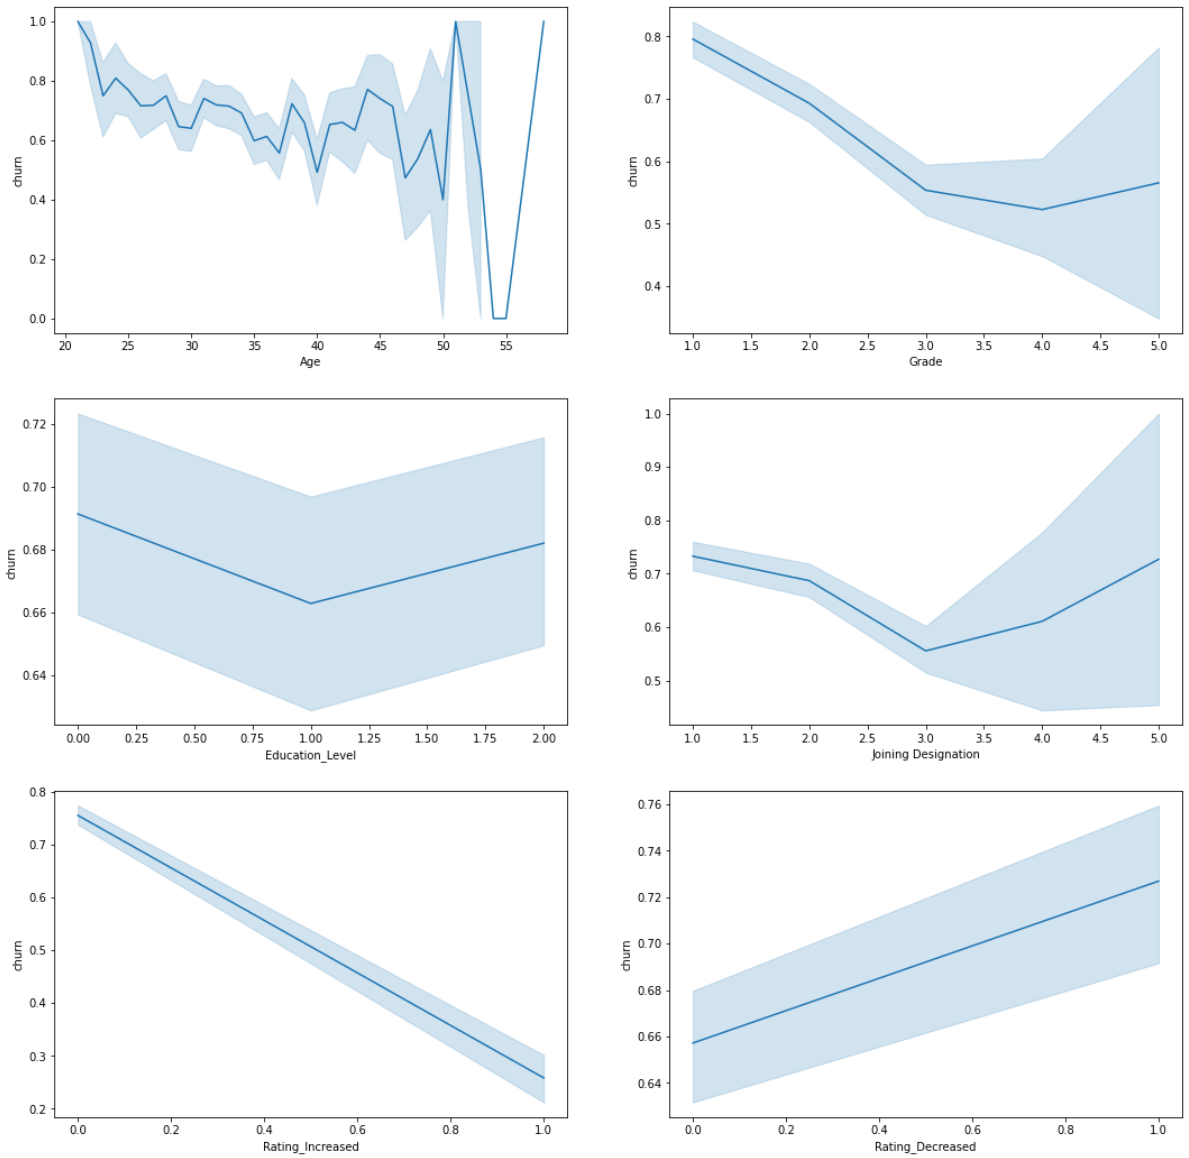
```

In [41]: fig, axis = plt.subplots(nrows=3, ncols=2, figsize=(18, 10))
fig.subplots_adjust(top=1.5)

sns.lineplot(x='Age', y='churn', data=df, ax=axis[0,0])
sns.lineplot(x='Grade', y='churn', data=df, ax=axis[0,1])
sns.lineplot(x='Education_Level', y='churn', data=df, ax=axis[1,0])
sns.lineplot(x='Joining Designation', y='churn', data=df, ax=axis[1,1])
sns.lineplot(x='Rating_Increased', y='churn', data=df, ax=axis[2,0])
sns.lineplot(x='Rating_Decreased', y='churn', data=df, ax=axis[2,1])

```




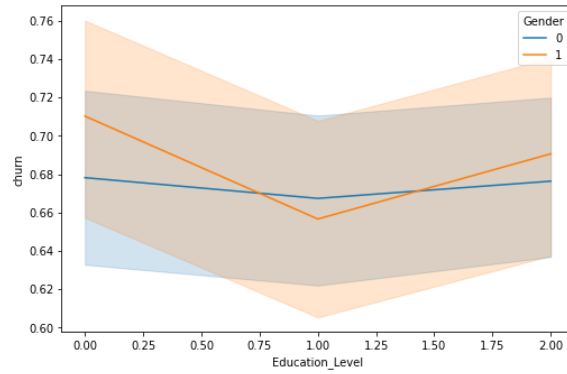


Out[41]: <AxesSubplot:xlabel='Rating\_Decreased', ylabel='churn'>



```
In [42]: fig, axis = plt.subplots(nrows=3, ncols=2, figsize=(18, 10))
fig.subplots_adjust(top=1.5)

sns.lineplot(x='Age', hue='Gender', y='churn', data=df, ax=axis[0,0])
sns.lineplot(x='Grade', hue='Gender', y='churn', data=df, ax=axis[0,1])
sns.lineplot(x='Education_Level', hue='Gender', y='churn', data=df, ax=axis[1,0])
sns.lineplot(x='Joining Designation', hue='Gender', y='churn', data=df, ax=axis[1,1])
sns.lineplot(x='Rating_Increased', hue='Gender', y='churn', data=df, ax=axis[2,0])
sns.lineplot(x='Rating_Deceased', hue='Gender', y='churn', data=df, ax=axis[2,1])
```

```
Out[42]: <AxesSubplot:xlabel='Rating_Deceased', ylabel='churn'>
```





```
In [43]: fig, axis = plt.subplots(nrows=2, figsize=(16, 6))
fig.subplots_adjust(top=1.9)

sns.lineplot(x='City',y='churn',data=df,ax=axis[0])
sns.lineplot(x='City',hue='Gender',y='churn',data=df,ax=axis[1])
```

```
Out[43]: <AxesSubplot:xlabel='City', ylabel='churn'>
```






In [44]: `## correlation matrix for heat map  
df.corr()`

Out[44]:

	Driver_ID	Age	Gender	Education_Level	Joining Designation	Grade	Income
Driver_ID	1.000000	-0.004829	0.015625	-0.014343	-0.023126	-0.014533	-0.017876
Age	-0.004829	1.000000	0.035064	-0.007876	0.079399	0.238908	0.209025
Gender	0.015625	0.035064	1.000000	-0.008747	-0.046815	-0.007765	0.009222
Education_Level	-0.014343	-0.007876	-0.008747	1.000000	0.003203	-0.016218	0.140189
Joining Designation	-0.023126	0.079399	-0.046815	0.003203	1.000000	0.726906	0.480523
Grade	-0.014533	0.238908	-0.007765	-0.016218	0.726906	1.000000	0.737210
Income	-0.017876	0.209025	0.009222	0.140189	0.480523	0.737210	1.000000
Total Business Value	0.015133	0.262775	0.018537	0.001392	-0.121368	0.340166	0.379426
Rating_Increased	0.009708	0.086854	0.036909	0.001734	0.015400	0.087604	0.110337
Rating_Decreased	0.002058	0.141208	-0.016209	0.061733	-0.172525	0.031182	0.098833
Quarterly Rating	0.012889	0.150336	0.021720	0.006544	-0.063404	0.085754	0.163426
churn	0.029269	-0.078571	0.012109	-0.007953	-0.127773	-0.205410	-0.201926

```
In [45]: sns.set(font_scale=1.15)
plt.figure(figsize=(15,8))
sns.heatmap(df.corr(),annot=True, vmin=-1, vmax=1)
```

Out[45]: <AxesSubplot:>




```
In [46]: ## spearman correlation matrix for heat map, used for better understanding
df.corr('spearman')
```

Out[46]:

	Driver_ID	Age	Gender	Education_Level	Joining Designation	Grade	Income
Driver_ID	1.000000	-0.007009	0.015667	-0.014416	-0.023266	-0.021009	-0.017635
Age	-0.007009	1.000000	0.033503	-0.012812	0.104681	0.250074	0.208455
Gender	0.015667	0.033503	1.000000	-0.008789	-0.042323	-0.009093	0.014659
Education_Level	-0.014416	-0.012812	-0.008789	1.000000	0.001773	-0.015532	0.144004
Joining Designation	-0.023266	0.104681	-0.042323	0.001773	1.000000	0.756124	0.482364
Grade	-0.021009	0.250074	-0.009093	-0.015532	0.756124	1.000000	0.706632
Income	-0.017635	0.208455	0.014659	0.144004	0.482364	0.706632	1.000000
Total Business Value	0.012676	0.254588	0.019311	0.036434	-0.167053	0.170581	0.247581
Rating_Increased	0.009885	0.085197	0.036909	0.001693	0.015131	0.094179	0.120519
Rating_Decreased	0.002057	0.127446	-0.016209	0.061635	-0.184552	0.008236	0.073446
Quarterly Rating	0.005634	0.147853	0.021989	0.005278	-0.057107	0.117197	0.178634
churn	0.029218	-0.083440	0.012109	-0.007874	-0.129816	-0.211135	-0.204218

```
In [47]: sns.set(font_scale=1.15)
plt.figure(figsize=(15,8))
sns.heatmap(df.corr('spearman'),annot=True,vmin=-1, vmax=1)
```

Out[47]: <AxesSubplot:>



## Observations:

- Rating Decreased is positively correlated with churn data irrespective of Gender.
- If age of driver is high there is high chance of that driver churns according lineplot.
- No impact of Grade and Income of driver churns.
- Gender is slightly impacting the churn data.
- Rating Increase is negatively correlated with churn irrespective of Gender.

## Data Preparation for Modeling

```
In [48]: df['City'].value_counts()
```

```
Out[48]: C20    152
          C15    101
          C29     96
          C26     93
          C8      89
          C27     89
          C10     86
          C16     84
          C22     82
          C3      82
          C28     82
          C12     81
          C5      80
          C1      80
          C21     79
          C14     79
          C6      78
          C4      77
          C7      76
          C9      75
          C25     74
          C23     74
          C24     73
          C19     72
          C2      72
          C17     71
          C13     71
          C18     69
          C11     64
          Name: City, dtype: int64
```

```
In [49]: one_hot_encoded_data = pd.get_dummies(df, columns = ['City'])
          one_hot_encoded_data
```

Out[49]:

	Driver_ID	MMM-YY	Age	Gender	Education_Level	Dateofjoining	LastWorkingDate	Join Designa
0	1	2019-03-01	28	0	2	2018-12-24	2019-03-11 00:00:00	
1	2	2020-12-01	31	0	2	2020-11-06	0	
2	4	2020-04-01	43	0	2	2019-12-07	2020-04-27 00:00:00	
3	5	2019-03-01	29	0	0	2019-01-09	2019-03-07 00:00:00	
4	6	2020-12-01	31	1	1	2020-07-31	0	
...	...	...	...	...	...	...	...	...
2376	2784	2020-12-01	34	0	0	2015-10-15	0	
2377	2785	2020-10-01	34	1	0	2020-08-28	2020-10-28 00:00:00	
2378	2786	2019-09-01	45	0	0	2018-07-31	2019-09-22 00:00:00	
2379	2787	2019-06-01	28	1	2	2018-07-21	2019-06-20 00:00:00	
2380	2788	2020-12-01	30	0	2	2020-06-08	0	

2381 rows × 44 columns

```

In [50]: df_numerical=one_hot_encoded_data.select_dtypes(exclude='object')
df_categorical=one_hot_encoded_data.select_dtypes(include='object')

In [53]: df_numerical['MMM-YY']=(df_numerical['MMM-YY'].astype(np.int64))
df_numerical['Dateofjoining']=(df_numerical['Dateofjoining'].astype(np.int64))

In [54]: def detect_outliers(data):
    length_before=len(data)
    q1=np.percentile(data,25)
    q3=np.percentile(data,75)
    IQR=q3-q1
    upper_bound=q3+1.5*IQR
    lower_bound=q1-1.5*IQR
    if lower_bound<0:
        lower_bound=0
    length_after=len(data[(data>lower_bound)&(data<upper_bound)])
    data=data[(data<=upper_bound) & (data>=lower_bound)]
    print('After applying IQR Method')
    return f"{np.round((length_before-length_after)/length_before,3)}% outliers da

In [55]: print("Before applying IQR Method:")
for col in df_numerical.columns:
    print(col,":", detect_outliers(df_numerical[col]))

```

Before applying IQR Method:  
After applying IQR Method  
Driver\_ID : 0.0% outliers data found  
After applying IQR Method  
MMM-YY : 0.0% outliers data found  
After applying IQR Method  
Age : 0.015% outliers data found  
After applying IQR Method  
Gender : 0.591% outliers data found  
After applying IQR Method  
Education\_Level : 0.329% outliers data found  
After applying IQR Method  
Dateofjoining : 0.06% outliers data found  
After applying IQR Method  
Joining Designation : 0.02% outliers data found  
After applying IQR Method  
Grade : 0.0% outliers data found  
After applying IQR Method  
Income : 0.02% outliers data found  
After applying IQR Method  
Total Business Value : 0.447% outliers data found  
After applying IQR Method  
Rating\_Increased : 1.0% outliers data found  
After applying IQR Method  
Rating\_Decreased : 0.691% outliers data found  
After applying IQR Method  
Quarterly Rating : 0.045% outliers data found  
After applying IQR Method  
churn : 0.321% outliers data found  
After applying IQR Method  
City\_C1 : 1.0% outliers data found  
After applying IQR Method  
City\_C10 : 1.0% outliers data found  
After applying IQR Method  
City\_C11 : 1.0% outliers data found  
After applying IQR Method  
City\_C12 : 1.0% outliers data found  
After applying IQR Method  
City\_C13 : 1.0% outliers data found  
After applying IQR Method  
City\_C14 : 1.0% outliers data found  
After applying IQR Method  
City\_C15 : 1.0% outliers data found  
After applying IQR Method  
City\_C16 : 1.0% outliers data found  
After applying IQR Method  
City\_C17 : 1.0% outliers data found  
After applying IQR Method  
City\_C18 : 1.0% outliers data found  
After applying IQR Method  
City\_C19 : 1.0% outliers data found  
After applying IQR Method  
City\_C2 : 1.0% outliers data found  
After applying IQR Method  
City\_C20 : 1.0% outliers data found  
After applying IQR Method  
City\_C21 : 1.0% outliers data found  
After applying IQR Method  
City\_C22 : 1.0% outliers data found  
After applying IQR Method  
City\_C23 : 1.0% outliers data found  
After applying IQR Method  
City\_C24 : 1.0% outliers data found  
After applying IQR Method

```

City_C25 : 1.0% outliers data found
After applying IQR Method
City_C26 : 1.0% outliers data found
After applying IQR Method
City_C27 : 1.0% outliers data found
After applying IQR Method
City_C28 : 1.0% outliers data found
After applying IQR Method
City_C29 : 1.0% outliers data found
After applying IQR Method
City_C3 : 1.0% outliers data found
After applying IQR Method
City_C4 : 1.0% outliers data found
After applying IQR Method
City_C5 : 1.0% outliers data found
After applying IQR Method
City_C6 : 1.0% outliers data found
After applying IQR Method
City_C7 : 1.0% outliers data found
After applying IQR Method
City_C8 : 1.0% outliers data found
After applying IQR Method
City_C9 : 1.0% outliers data found

```

## Observations:

- Very less amount of data is present and outliers are too very less,
- lets not remove outliers and use this data to train the model.

## Ensemble -Bagging Model

```

In [57]: targets = df_numerical['churn']
         features = df_numerical.drop(columns = {'churn'})

In [59]: # splitting the data into train and test with some test size
         from sklearn.model_selection import train_test_split
         X_train, X_test, y_train, y_test = train_test_split(features, targets, test_size=0

In [60]: # scaling the X_train and X_test
         from sklearn.preprocessing import StandardScaler
         scaler = StandardScaler()

         X_train = scaler.fit_transform(X_train)
         X_test = scaler.fit_transform(X_test)

In [71]: #Building Ensemble Bagging Model
         from sklearn.ensemble import RandomForestClassifier

         model = RandomForestClassifier(random_state=7, max_depth=3, n_estimators=50)

         #training the Model
         model.fit(X_train,y_train)

         #accuracy score
         accuracy=model.score(X_test, y_test)

         print(f'accuracy score: {accuracy}')

```

accuracy score: 0.9161425576519916

```
In [62]: # Handling Imbalance data with SMOTE Technique
from imblearn.over_sampling import SMOTE
from collections import Counter

smt = SMOTE()
X_sm, y_sm = smt.fit_resample(X_train, y_train)

print('Resampled dataset shape {}'.format(Counter(y_sm)))
```

Resampled dataset shape Counter({0: 1289, 1: 1289})

```
In [63]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold, cross_validate

tree_clf = RandomForestClassifier(random_state=7, max_depth=3, n_estimators=50)
kfold = KFold(n_splits=10)
cv_acc_results = cross_validate(tree_clf, X_sm, y_sm, cv = kfold, scoring = 'accuracy')


print(f"K-Fold Accuracy Mean: Train: {cv_acc_results['train_score'].mean()*100} Validation: {cv_acc_results['test_score'].mean()*100}")
print(f"K-Fold Accuracy Std: Train: {cv_acc_results['train_score'].std()*100} Validation: {cv_acc_results['test_score'].std()*100}")
```

K-Fold Accuracy Mean: Train: 94.75474676491999 Validation: 93.9094501251772

K-Fold Accuracy Std: Train: 0.22269082954512817 Validation: 1.6124656378346365

```
In [64]: # Feature importance
import numpy as np
import matplotlib.pyplot as plt

clf = RandomForestClassifier(random_state=7, max_depth=3, n_estimators=50)
clf.fit(X_train, y_train)
importances = clf.feature_importances_
indices = np.argsort(importances)[::-1] # Sort feature importances in descending order
names = [features.columns[i] for i in indices] # Rearrange feature names so they match indices
plt.figure(figsize=(15, 7)) # Create plot
plt.title("Feature Importance") # Create plot title
plt.bar(range(features.shape[1]), importances[indices]) # Add bars
plt.xticks(range(features.shape[1]), names, rotation=90) # Add feature names as x-axis labels
plt.show() # Show plot
```





# Observations:

- we got Accuracy of 92% using Random Forest Algorithm.
- Let's try Hyper parameter Tuning to get best Random Forest Parameters.

## HyperParameter Tuning

```
In [65]: # Defining Parametes
params = {
    'n_estimators' : [20,30,40,50],
    'max_depth' : [1,2,3],
    'criterion' : ['gini', 'entropy'],
    'bootstrap' : [True, False],
    'max_features' : [2,3,4,5,6,7],
}
```

```
In [72]: # Using GridSearchCV to get best Parameters for Model.
from sklearn.model_selection import GridSearchCV

# Tuning Function
tuning_function = GridSearchCV(estimator = RandomForestClassifier(),
                               param_grid = params,
                               scoring = 'accuracy',
                               cv = 3,
                               n_jobs=-1
                              )

# Fitting the Tuning Function
tuning_function.fit(X_train, y_train)

# Fetching best parameters and score
parameters = tuning_function.best_params_
score = tuning_function.best_score_
print(parameters)
print(score)

{'bootstrap': True, 'criterion': 'entropy', 'max_depth': 3, 'max_features': 7, 'n_
estimators': 20}
0.9574612715997252
```

```
In [80]: #Lets train the model with the best parameters.
#Building Ensemble Bagging Model
from sklearn.ensemble import RandomForestClassifier

model = RandomForestClassifier(random_state=7, max_depth=3, n_estimators=20, max_f

#training the Model
model.fit(X_train, y_train)

#Predicting the model
y_pred=model.predict(X_test)

#accuracy score
accuracy=model.score(X_test, y_test)

print(f'accuracy score: {accuracy}')
```

accuracy score: 0.9559748427672956

```
In [74]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold, cross_validate


tree_clf = RandomForestClassifier(random_state=7, max_depth=3, n_estimators=20, max
kfold = KFold(n_splits=10)
cv_acc_results = cross_validate(tree_clf, X_sm, y_sm, cv = kfold, scoring = 'accuracy')

print(f"K-Fold Accuracy Mean: Train: {cv_acc_results['train_score'].mean()*100} Validation: {cv_acc_results['test_score'].mean()*100}")
print(f"K-Fold Accuracy Std: Train: {cv_acc_results['train_score'].std()*100} Validation: {cv_acc_results['test_score'].std()*100}")
```

K-Fold Accuracy Mean: Train: 94.98747381479447 Validation: 94.6852471872832  
K-Fold Accuracy Std: Train: 0.258429506253071 Validation: 1.3008339018818098

```
In [75]: # Feature importance
import numpy as np
import matplotlib.pyplot as plt

clf = RandomForestClassifier(random_state=7, max_depth=3, n_estimators=50)
clf.fit(X_train, y_train)
importances = clf.feature_importances_
indices = np.argsort(importances)[::-1] # Sort feature importances in descending order
names = [features.columns[i] for i in indices] # Rearrange feature names so they match indices
plt.figure(figsize=(15, 7)) # Create plot
plt.title("Feature Importance") # Create plot title
plt.bar(range(features.shape[1]), importances[indices]) # Add bars
plt.xticks(range(features.shape[1]), names, rotation=90) # Add feature names as x-axis labels
plt.show() # Show plot
```




## Observations:

- after applying Hyperparameter Tuning we got 95.5% Accuracy.

```
In [81]: # Confusion Matrix
from sklearn.metrics import confusion_matrix, precision_score, recall_score, ConfusionMatrixDisplay
cm = confusion_matrix(y_test, y_pred)
cm
```

```
Out[81]: array([[141,  9],
                [ 12, 315]], dtype=int64)
```

```
In [82]: disp=ConfusionMatrixDisplay(confusion_matrix=cm,
...                                  display_labels=model.classes_)
disp.plot()
plt.show()
```



- TN=True Negatives
- TP=True Positives
- FN=False Negatives
- FP=False Positives

```
In [83]: TN=cm[0][0]
TP=cm[1][1]
FP=cm[0][1]
FN=cm[1][0]

print('True Negatives',TN)
print('True Positives',TP)
print('False Positives',FP)
print('False Negatives',FN)
```

```
True Negatives 141
True Positives 315
False Positives 9
False Negatives 12
```

```
In [84]: # calculating Accuracy
def Accuracy(TN,TP,FP,FN):
    ans=(TN+TP)/(TN+TP+FP+FN)
    return ans
```

```
In [85]: Accuracy(TN,TP,FP,FN)
```

```
Out[85]: 0.9559748427672956
```

```
In [86]: Precision=precision_score(y_test,y_pred)
Precision
```

```
Out[86]: 0.9722222222222222
```

```
In [87]: Recall=recall_score(y_test,y_pred)
Recall
```

```
Out[87]: 0.963302752293578
```

```
In [88]: # Calculating F1 Score
def F1_score(Precision,Recall):
    ans=2*(Precision*Recall)/(Precision+Recall)
    return ans
```

```
In [89]: F1_score=F1_score(Precision,Recall)
F1_score
```


```
Out[89]: 0.9677419354838711
```

```
In [90]: probs=model.predict_proba(X_test)[: ,1]
```

```
In [91]: # ROC AUC Curve
from sklearn.metrics import roc_auc_score,roc_curve
fpr,tpr,thres=roc_curve(y_test,probs)
plt.plot(fpr,tpr)
print('roc_auc_score:',roc_auc_score(y_test,probs))
# axis Labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

```
roc_auc_score: 0.9613761467889909
```

```
Out[91]: Text(0, 0.5, 'True Positive Rate')
```



## Observations:

- Here, AUC score is greater than 0.50. it means our model can predict the positive & Negative class well.
- Max score can be 1 but we got 0.96.
- Our Model is doing Great according to this Area under curve

```
In [92]: # Precision-Recall Curve
from sklearn.metrics import auc,precision_recall_curve,f1_score
lr_precision, lr_recall, _ = precision_recall_curve(y_test,y_pred)
lr_f1 = f1_score(y_test, y_pred)
```


```

auc(lr_recall,lr_precision)
print("f1 score",lr_f1)
plt.plot(lr_recall,lr_precision,marker='.')
# axis labels
plt.xlabel('Recall')
plt.ylabel('Precision')

```

f1 score 0.9677419354838711

Out[92]: Text(0, 0.5, 'Precision')



## Observations:

- F1 score is 0.967 which is very high
- According to this curve we get that There are high Precision and High Recall.
- This implies there are less false positives and false negatives.

## Ensemble - Boosting Method

```

In [94]: from xgboost import XGBClassifier
from sklearn.model_selection import RandomizedSearchCV,GridSearchCV
from sklearn.model_selection import StratifiedKFold
import datetime as dt

```

```

In [97]: #Building Ensemble Boosting Model
from sklearn.ensemble import RandomForestClassifier

xgb_model = XGBClassifier(n_estimators=40, objective='multi:softmax', num_class=15)

#training the Model
xgb_model.fit(X_train,y_train)

#accuracy score
accuracy=xgb_model.score(X_test, y_test)

print(f'accuracy score: {accuracy}')

```

[20:00:56] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

accuracy score: 0.9538784067085954

```
In [98]: # Handling Imbalance data with SMOTE Technique
from imblearn.over_sampling import SMOTE
from collections import Counter

smt = SMOTE()
X_sm, y_sm = smt.fit_resample(X_train, y_train)

print('Resampled dataset shape {}'.format(Counter(y_sm)))
```

Resampled dataset shape Counter({0: 1289, 1: 1289})

```
In [106... from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold, cross_validate

xgb_model = XGBClassifier(n_estimators=40, objective='multi:softmax', num_class=15)
kfold = KFold(n_splits=10)
cv_acc_results = cross_validate(xgb_model, X_sm, y_sm, cv = kfold, scoring = 'accuracy')

print(f"K-Fold Accuracy Mean: Train: {cv_acc_results['train_score'].mean()*100} Validation: {cv_acc_results['val_score'].mean()*100}")
print(f"K-Fold Accuracy Std: Train: {cv_acc_results['train_score'].std()*100} Validation: {cv_acc_results['val_score'].std()*100}")
```

[20:04:42] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:42] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:42] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:42] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:43] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:43] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:43] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:44] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:44] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:04:44] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.


K-Fold Accuracy Mean: Train: 99.91380424608893 Validation: 95.73371942207342  
K-Fold Accuracy Std: Train: 0.050988135452225684 Validation: 1.5000760375190425

In [107...

```
# Feature importance
import numpy as np
import matplotlib.pyplot as plt

xgb_model = XGBClassifier(n_estimators=40, objective='multi:softmax', num_class=15)
xgb_model.fit(X_train, y_train)
importances = clf.feature_importances_
indices = np.argsort(importances)[::-1] # Sort feature importances in descending order
names = [features.columns[i] for i in indices] # Rearrange feature names so they match indices
plt.figure(figsize=(15, 7)) # Create plot
plt.title("Feature Importance") # Create plot title
plt.bar(range(features.shape[1]), importances[indices]) # Add bars
plt.xticks(range(features.shape[1]), names, rotation=90) # Add feature names as x-axis labels
plt.show() # Show plot
```

[20:05:49] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.



## Observations:

- we got Accuracy of 95.3% using XGBoost Algorithm.
- Let's try Hyper parameter Tuning to get best XGBoost Parameters.

## HyperParameter Tuning

```
In [108... # Defining Parametes
params = {
    'n_estimators' : [20,30,40,50],
    'max_depth' : [1,2,3],
    'learning_rate': [0.1, 0.5, 0.8],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
}
```

```
In [ ]: folds = 3

skf = StratifiedKFold(n_splits=folds, shuffle = True, random_state = 1001)

random_search = RandomizedSearchCV(xgb, param_distributions=params, n_iter=10, sco

random_search.fit(X_sm, y_sm)
print(random_search.best_params_)
```

```
In [109... # Using GridSearchCV to get best Parameters for Model.
from sklearn.model_selection import GridSearchCV

# Tuning Function
tuning_function = GridSearchCV(estimator = XGBClassifier(),
                               param_grid = params,
                               scoring = 'accuracy',
                               cv = 3,
                               n_jobs=-1)
```



```

    )

# Fitting the Tuning Function
tuning_function.fit(X_train, y_train)

# Fetching best parameters and score
parameters = tuning_function.best_params_
score = tuning_function.best_score_
print(parameters)
print(score)

{'colsample_bytree': 1.0, 'learning_rate': 0.8, 'max_depth': 1, 'n_estimators': 3
0, 'subsample': 1.0}
0.9574629275441516

```

```

In [111... #Building Ensemble Boosting Model
from sklearn.ensemble import RandomForestClassifier

xgb_model = XGBClassifier(n_estimators=40, objective='multi:softmax', num_class=15,
                          max_depth=1, subsample=1.0, silent=True)

#training the Model
xgb_model.fit(X_train, y_train)

#accuracy score
accuracy=xgb_model.score(X_test, y_test)

print(f'accuracy score: {accuracy}')

[20:12:21] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-gr
oup-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:
Parameters: { "silent" } are not used.

accuracy score: 0.9559748427672956

```

```

In [112... from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold, cross_validate

xgb_model = XGBClassifier(n_estimators=40, objective='multi:softmax', num_class=15,
                          max_depth=1, subsample=1.0, silent=True)
kfold = KFold(n_splits=10)
cv_acc_results = cross_validate(xgb_model, X_sm, y_sm, cv = kfold, scoring = 'accu

print(f"K-Fold Accuracy Mean: Train: {cv_acc_results['train_score'].mean()*100} Val
print(f"K-Fold Accuracy Std: Train: {cv_acc_results['train_score'].std()*100} Valid

```

[20:13:28] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:29] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:29] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:29] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:29] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:29] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:29] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:30] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:30] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.

[20:13:30] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.


K-Fold Accuracy Mean: Train: 96.21586452034646 Validation: 95.69465810032275  
K-Fold Accuracy Std: Train: 0.2523114097029877 Validation: 1.1152665125941432

In [113...

```
# Feature importance
import numpy as np
import matplotlib.pyplot as plt

xgb_model = XGBClassifier(n_estimators=40, objective='multi:softmax', num_class=15,
                          max_depth=1, subsample=1.0, silent=True)
xgb_model.fit(X_train, y_train)
importances = clf.feature_importances_
indices = np.argsort(importances)[::-1] # Sort feature importances in descending order
names = [features.columns[i] for i in indices] # Rearrange feature names so they match indices
plt.figure(figsize=(15, 7)) # Create plot
plt.title("Feature Importance") # Create plot title
plt.bar(range(features.shape[1]), importances[indices]) # Add bars
plt.xticks(range(features.shape[1]), names, rotation=90) # Add feature names as x-axis labels
plt.show() # Show plot
```

[20:14:05] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-0fc7796c793e6356f-1/xgboost/xgboost-ci-windows/src/learner.cc:767:  
Parameters: { "silent" } are not used.




## Observations:

- After Applying Hyperparameter Tuning we got 95.5% Accuracy.

```
In [114... # Confusion Matrix
from sklearn.metrics import confusion_matrix, precision_score, recall_score, ConfusionMatrixDisplay
cm = confusion_matrix(y_test, y_pred)
cm
```

```
Out[114]: array([[141,  9],
       [ 13, 314]], dtype=int64)
```

```
In [115... disp=ConfusionMatrixDisplay(confusion_matrix=cm,
                               display_labels=model.classes_)
disp.plot()
plt.show()
```



```
In [116... TN=cm[0][0]
TP=cm[1][1]
FP=cm[0][1]
```

```
FN=cm[1][0]
```

```
print('True Negatives',TN)
print('True Positives',TP)
print('False Positives',FP)
print('False Negatives',FN)
```

```
True Negatives 141
True Positives 314
False Positives 9
False Negatives 13
```

```
In [117... Accuracy(TN,TP,FP,FN)
```

```
Out[117]: 0.9538784067085954
```

```
In [118... Precision=precision_score(y_test,y_pred)
Precision
```

```
Out[118]: 0.9721362229102167
```

```
In [119... Recall=recall_score(y_test,y_pred)
Recall
```

```
Out[119]: 0.9602446483180428
```

```
In [126... # Calculating F1 Score
def F1_score(Precision,Recall):
    ans=2*(Precision*Recall)/(Precision+Recall)
    return ans
```


```
In [127... F1_score=F1_score(Precision,Recall)
F1_score
```

```
Out[127]: 0.9721362229102167
```

```
In [128... probs=model.predict_proba(X_test)[: ,1]
```

```
In [129... # ROC AUC Curve
from sklearn.metrics import roc_auc_score,roc_curve
fpr,tpr,thres=roc_curve(y_test,probs)
plt.plot(fpr,tpr)
print('roc_auc_score:',roc_auc_score(y_test,probs))
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

```
roc_auc_score: 0.9613761467889909
Out[129]: Text(0, 0.5, 'True Positive Rate')
```




## Observations:

- Here, AUC score is greater than 0.50. it means our model can predict the positive & Negative class well.
- Max score can be 1 but we got 0.96.
- Our Model is doing Great according to this Area under curve

```
In [130... # Precision-Recall Curve
from sklearn.metrics import auc, precision_recall_curve, f1_score
lr_precision, lr_recall, _ = precision_recall_curve(y_test, y_pred)
lr_f1 = f1_score(y_test, y_pred)
auc(lr_recall, lr_precision)
print("f1 score", lr_f1)
plt.plot(lr_recall, lr_precision, marker='.')
# axis Labels
plt.xlabel('Recall')
plt.ylabel('Precision')
```

f1 score 0.9661538461538461  
Text(0, 0.5, 'Precision')

Out[130]:



## Observations:

- F1 score is 0.966 which is very high
- According to this curve we get that There are high Precision and High Recall.
- This implies there are less false positives and false negatives.

## Actionable insights:

- No Duplicate Records Found
- some missing and outliers are present later imputed them.
- No change in monthly income so not including that column for further analysis.
- Major Drivers are male according to Dataset.
- Total of 368 drivers rating has been increased.
- Total of 973 drivers rating has been Decreased.
- Major of Drivers are Graduated according to Dataset.
- More drivers are from city c20.
- The average age for Drivers are 30-35.
- Rating Decreased is positively correlated with churn data irrespective of Gender.
- If age of driver is high there is high chance of that driver churns according lineplot.
- No impact of Grade and Income of driver churns.
- Gender is slightly impacting the churn data.
- Rating Increase is negatively correlated with churn irrespective of Gender.
- F1-score for Bagging(Random Forest) is 96.77
- F1-score for Boosting(XGBoost) is 96.61

## Recommendations:

- Drivers with Quarterly Rating Increased likely not to churn so provide some incentives or bonus to such Drivers.
- Drivers with Quarterly Rating Reduced likely to churn more, so try to connect with customers regarding the issues faced and try to clear that issues.
- Provide some paid leaves for Drivers to reduce churning of Drivers.
- Take Surveys in detailed manner to know the issues faced by Drivers.
- Payment might be the major reason try to monitor the payments paid to Driver.

In [ ]: